

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from xgboost import XGBClassifier
from sklearn import metrics

import warnings
warnings.filterwarnings('ignore')
```

```
df = pd.read_csv('/content/TSLA.csv')
df.head()
```

	Date	Open	High	Low	Close	Adj Close	Volume
0	2010-06-29	19.000000	25.00	17.540001	23.889999	23.889999	18766300
1	2010-06-30	25.790001	30.42	23.299999	23.830000	23.830000	17187100
2	2010-07-01	25.000000	25.92	20.270000	21.959999	21.959999	8218800
3	2010-07-02	23.000000	23.10	18.709999	19.200001	19.200001	5139800
4	2010-07-06	20.000000	20.00	15.830000	16.110001	16.110001	6866900

```
df.shape
```

```
(2416, 7)
```

```
df.describe()
```

	Open	High	Low	Close	Adj Close	Volume
count	2416.000000	2416.000000	2416.000000	2416.000000	2416.000000	2.416000e+03
mean	186.271147	189.578224	182.916639	186.403651	186.403651	5.572722e+06
std	118.740163	120.892329	116.857591	119.136020	119.136020	4.987809e+06
min	16.139999	16.629999	14.980000	15.800000	15.800000	1.185000e+05
25%	34.342498	34.897501	33.587501	34.400002	34.400002	1.899275e+06
50%	213.035004	216.745002	208.870002	212.960007	212.960007	4.578400e+06
75%	266.450012	270.927513	262.102501	266.774994	266.774994	7.361150e+06
max	673.690002	786.140015	673.520020	780.000000	780.000000	4.706500e+07

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2416 entries, 0 to 2415
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Date         2416 non-null   object
1   Open         2416 non-null   float64
2   High         2416 non-null   float64
3   Low          2416 non-null   float64
4   Close        2416 non-null   float64
5   Adj Close    2416 non-null   float64
6   Volume       2416 non-null   int64
dtypes: float64(5), int64(1), object(1)
memory usage: 132.2+ KB
```

```
plt.figure(figsize=(15,5))
plt.plot(df['Close'])
plt.title('Tesla Close price.', fontsize=15)
plt.ylabel('Price in dollars.')
plt.show()
```



```
df.head()
```

	Date	Open	High	Low	Close	Adj Close	Volume
0	2010-06-29	19.000000	25.00	17.540001	23.889999	23.889999	18766300
1	2010-06-30	25.790001	30.42	23.299999	23.830000	23.830000	17187100
2	2010-07-01	25.000000	25.92	20.270000	21.959999	21.959999	8218800
3	2010-07-02	23.000000	23.10	18.709999	19.200001	19.200001	5139800
4	2010-07-06	20.000000	20.00	15.830000	16.110001	16.110001	6866900

```
df[df['Close'] == df['Adj Close']].shape
```

```
(2416, 7)
```

```
df = df.drop(['Adj Close'], axis=1)
```

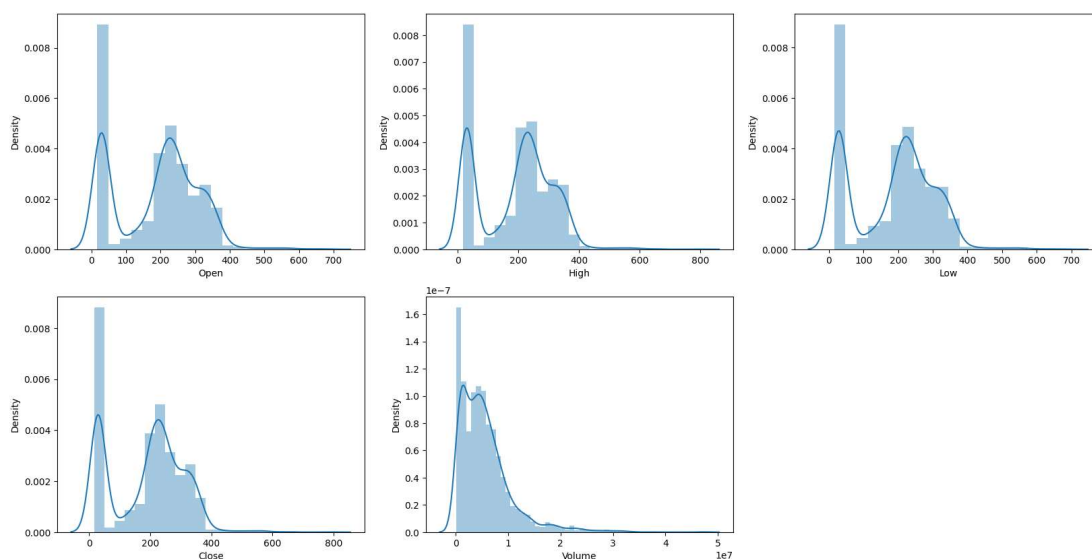
```
df.isnull().sum()
```

```
Date      0
Open      0
High      0
Low       0
Close     0
Volume    0
dtype: int64
```

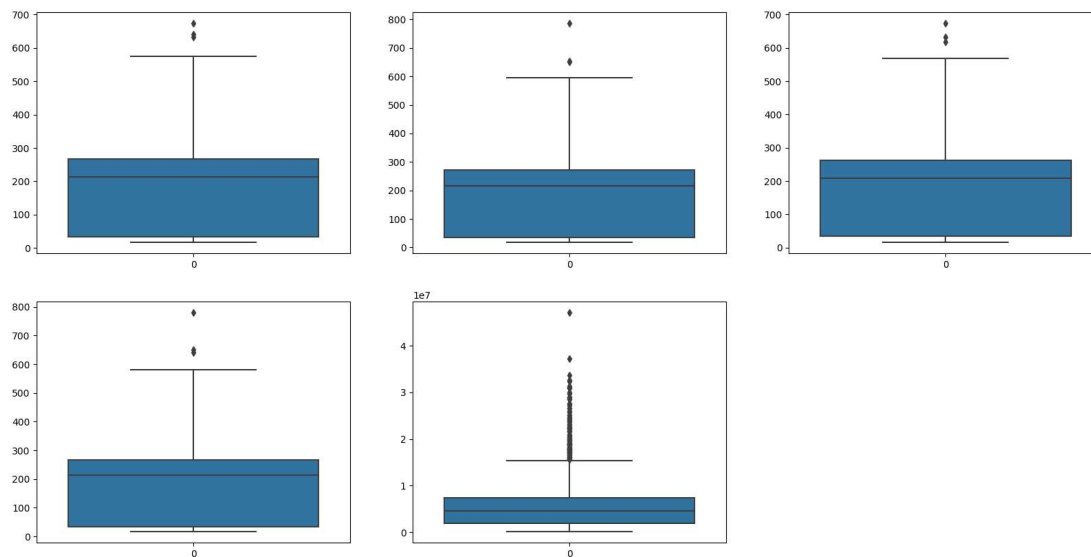
```
features = ['Open', 'High', 'Low', 'Close', 'Volume']
```

```
plt.subplots(figsize=(20,10))
```

```
for i, col in enumerate(features):
    plt.subplot(2,3,i+1)
    sb.distplot(df[col])
plt.show()
```



```
plt.subplots(figsize=(20,10))
for i, col in enumerate(features):
    plt.subplot(2,3,i+1)
    sb.boxplot(df[col])
plt.show()
```



```
splitted = df['Date'].str.split('-', expand=True)
```

```
df['day'] = splitted[1].astype('int')
df['month'] = splitted[0].astype('int')
df['year'] = splitted[2].astype('int')
```

```
df.head()
```

	Date	Open	High	Low	Close	Volume	day	month	year
0	2010-06-29	19.000000	25.00	17.540001	23.889999	18766300	6	2010	29
1	2010-06-30	25.790001	30.42	23.299999	23.830000	17187100	6	2010	30
2	2010-07-01	25.000000	25.92	20.270000	21.959999	8218800	7	2010	1
3	2010-07-02	23.000000	23.10	18.709999	19.200001	5139800	7	2010	2
4	2010-07-06	20.000000	20.00	15.830000	16.110001	6866900	7	2010	6





```
df['is_quarter_end'] = np.where(df['month']%3==0,1,0)
df.head()
```

	Date	Open	High	Low	Close	Volume	day	month	year	is_quarter_e
0	2010-06-29	19.000000	25.00	17.540001	23.889999	18766300	6	2010	29	
1	2010-06-30	25.790001	30.42	23.299999	23.830000	17187100	6	2010	30	
2	2010-07-01	25.000000	25.92	20.270000	21.959999	8218800	7	2010	1	

```
data_grouped = df.groupby('year').mean()
plt.subplots(figsize=(20,10))
```

```
for i, col in enumerate(['Open', 'High', 'Low', 'Close']):
    plt.subplot(2,2,i+1)
    data_grouped[col].plot.bar()
```

```
plt.show()
```




```
df.groupby('is_quarter_end').mean()
```

	Open	High	Low	Close	Volume	day
is_quarter_end						
0	195.383850	198.830719	191.934386	195.566477	5.012534e+06	6.455556 20
1	170.531763	173.600440	167.341233	170.580711	6.540089e+06	6.972012 20

```
df['open-close'] = df['Open'] - df['Close']
df['low-high'] = df['Low'] - df['High']
df['target'] = np.where(df['Close'].shift(-1) > df['Close'], 1, 0)
```

```
plt.pie(df['target'].value_counts().values,
        labels=[0, 1], autopct='%1.1f%%')
plt.show()
```



```
plt.figure(figsize=(10, 10))
```

```
# As our concern is with the highly  
# correlated features only so, we will visualize  
# our heatmap as per that criteria only.  
sb.heatmap(df.corr() > 0.9, annot=True, cbar=False)  
plt.show()
```

Open	1	1	1	1	0	0	0	0	0	0	0	0
High	1	1	1	1	0	0	0	0	0	0	0	0

```
features = df[['open-close', 'low-high', 'is_quarter_end']]
target = df['target']
```

```
scaler = StandardScaler()
features = scaler.fit_transform(features)
```

```
X_train, X_valid, Y_train, Y_valid = train_test_split(
    features, target, test_size=0.1, random_state=2022)
print(X_train.shape, X_valid.shape)
```

```
(2174, 3) (242, 3)
```

```
models = [LogisticRegression(), SVC(
    kernel='poly', probability=True), XGBClassifier()]
```

```
for i in range(3):
    models[i].fit(X_train, Y_train)

    print(f'{models[i]} : ')
    print('Training Accuracy : ', metrics.roc_auc_score(
        Y_train, models[i].predict_proba(X_train)[:,-1]))
    print('Validation Accuracy : ', metrics.roc_auc_score(
        Y_valid, models[i].predict_proba(X_valid)[:,-1]))
    print()
```

```
LogisticRegression() :
Training Accuracy : 0.5089337346592329
Validation Accuracy : 0.5124521072796935
```

```
SVC(kernel='poly', probability=True) :
Training Accuracy : 0.5165199795103654
Validation Accuracy : 0.4817323481116584
```

```
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric=None, feature_types=None,
               gamma=None, grow_policy=None, importance_type=None,
               interaction_constraints=None, learning_rate=None, max_bin=None,
               max_cat_threshold=None, max_cat_to_onehot=None,
               max_delta_step=None, max_depth=None, max_leaves=None,
               min_child_weight=None, missing=nan, monotone_constraints=None,
               multi_strategy=None, n_estimators=None, n_jobs=None,
               num_parallel_tree=None, random_state=None, ...) :
Training Accuracy : 0.9365469039061541
Validation Accuracy : 0.45583607006020804
```

```
metrics.plot_confusion_matrix(models[0], X_valid, Y_valid)
plt.show()
```

```
-----
AttributeError                                Traceback (most recent call last)
<ipython-input-34-e4a3478fce11> in <cell line: 1>()
----> 1 metrics.plot_confusion_matrix(models[0], X_valid, Y_valid)
      2 plt.show()
```

```
"""AttributeError: module 'sklearn.metrics' has no attribute 'plot_confusion_matrix'
it is been deprecated from sklearn"""
```

```
'AttributeError: module 'sklearn.metrics' has no attribute 'plot_confusion_matrix'\nIt
is been deprecated from sklearn'
```

```
"""Conclusion:
```

```
We can observe that the accuracy achieved by the state-of-the-art ML model is no better than simply guessing with a
```

```
I have tried to make it more better """
```

```
'Conclusion:\nWe can observe that the accuracy achieved by the state-of-the-art ML mode
l is no better than simply guessing with a probability of 50%. Possible reasons for thi
s may be the lack of data or using a very simple model to perform such a complex task a
s Stock Market prediction \n\nI have tried to make it more better '
```

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from xgboost import XGBClassifier
from sklearn import metrics

# Load your dataset
df = pd.read_csv('/content/TSLA.csv')

# Data Preprocessing
# ... (You can include feature engineering, data cleaning, etc.)

# Feature Scaling
scaler = StandardScaler()
features = scaler.fit_transform(features)

# Data Splitting
X_train, X_valid, Y_train, Y_valid = train_test_split(features, target, test_size=0.1, random_state=2022)

# Model Development
model = XGBClassifier(n_estimators=100, max_depth=3, learning_rate=0.1)

model.fit(X_train, Y_train)

# Evaluation
train_accuracy = metrics.roc_auc_score(Y_train, model.predict_proba(X_train)[: , 1])
valid_accuracy = metrics.roc_auc_score(Y_valid, model.predict_proba(X_valid)[: , 1])

print(f'Training Accuracy: {train_accuracy}')
print(f'Validation Accuracy: {valid_accuracy}')

Training Accuracy: 0.710399759543132
Validation Accuracy: 0.504447181171319

df.head(5)
```


	Date	Open	High	Low	Close	Volume	open-close	low-high	target
0	2010-06-29	19.000000	25.00	17.540001	23.889999	18766300	-4.889999	-7.459999	0
1	2010-06-30	25.790001	30.42	23.299999	23.830000	17187100	1.960001	-7.120001	0
2	2010-	25.000000	25.92	20.270000	21.959999	8218800	3.040001	-5.650000	0

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from xgboost import XGBClassifier
from sklearn import metrics
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.model_selection import GridSearchCV
```

```
# Load your dataset, perform initial data preprocessing, and feature engineering
# Replace this with your data loading and preprocessing steps
```

```
df = pd.read_csv('/content/TSLA.csv')
df = df.drop(['Adj Close'], axis=1)
```

```
# Feature engineering
df['open-close'] = df['Open'] - df['Close']
df['low-high'] = df['Low'] - df['High']
df['target'] = np.where(df['Close'].shift(-1) > df['Close'], 1, 0)
```

```
# Split the data into features and target
features = df[['open-close', 'low-high']]
target = df['target']
```

```
# Normalize the features
scaler = StandardScaler()
features = scaler.fit_transform(features)
```

```
# Split the data into training and validation sets
X_train, X_valid, Y_train, Y_valid = train_test_split(
    features, target, test_size=0.1, random_state=2022)
```

```
# Hyperparameter tuning using GridSearchCV
param_grid = {
    'max_depth': [3, 5, 7],
    'n_estimators': [50, 100, 200],
    'learning_rate': [0.01, 0.1, 0.2]
}
```

```
xgb = XGBClassifier()
grid_search = GridSearchCV(xgb, param_grid, cv=3)
grid_search.fit(X_train, Y_train)
best_xgb = grid_search.best_estimator_
```

```
# Ensemble multiple models
models = [best_xgb, RandomForestClassifier(), GradientBoostingClassifier()]
```

```
for model in models:
    model.fit(X_train, Y_train)
    print(f'{model} : ')
    train_accuracy = metrics.roc_auc_score(Y_train, model.predict_proba(X_train)[:, 1])
    valid_accuracy = metrics.roc_auc_score(Y_valid, model.predict_proba(X_valid)[:, 1])
```

```
print(f'Training Accuracy: {train_accuracy}')
print(f'Validation Accuracy: {valid_accuracy}')
print()
```

```
# Combine predictions from multiple models using majority voting or averaging
```

```
# Evaluate the ensemble model
```

```
# Continue optimizing and experimenting with different techniques to improve accuracy
```

```
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric=None, feature_types=None,
               gamma=None, grow_policy=None, importance_type=None,
               interaction_constraints=None, learning_rate=0.01, max_bin=None,
               max_cat_threshold=None, max_cat_to_onehot=None,
               max_delta_step=None, max_depth=5, max_leaves=None,
               min_child_weight=None, missing=nan, monotone_constraints=None,
               multi_strategy=None, n_estimators=50, n_jobs=None,
               num_parallel_tree=None, random_state=None, ...) :
```

```
Training Accuracy: 0.6684091322809111
```

```
Validation Accuracy: 0.45740968801313625
```

```
RandomForestClassifier() :
```

```
Training Accuracy: 0.9999995766604435
```

```
Validation Accuracy: 0.43045292829775594
```

```
GradientBoostingClassifier() :
```

```
Training Accuracy: 0.772182781086882
```

```
Validation Accuracy: 0.4714354132457581
```

```
from sklearn.metrics import plot_confusion_matrix
import matplotlib.pyplot as plt
```

```
# Replace 'best_xgb' with your chosen model
```

```
plot_confusion_matrix(best_xgb, X_valid, Y_valid)
```

```
plt.show()
```

```
-----
ImportError                                Traceback (most recent call last)
```

```
<ipython-input-49-c48be53b5521> in <cell line: 1>()
```

```
----> 1 from sklearn.metrics import plot_confusion_matrix
```

```
      2 import matplotlib.pyplot as plt
```

```
      3
```

```
      4 # Replace 'best_xgb' with your chosen model
```

```
      5 plot_confusion_matrix(best_xgb, X_valid, Y_valid)
```

```
ImportError: cannot import name 'plot_confusion_matrix' from 'sklearn.metrics'
(/usr/local/lib/python3.10/dist-packages/sklearn/metrics/__init__.py)
```

```
-----
NOTE: If your import is failing due to a missing package, you can
manually install dependencies using either !pip or !apt.
```

```
To view examples of installing some common dependencies, click the
"Open Examples" button below.
```

OPEN EXAMPLES

SEARCH STACK OVERFLOW

```
pip install -U scikit-learn
```

```
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (
Collecting scikit-learn
```

```
  Downloading scikit_learn-1.3.2-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.
    10.8/10.8 MB 77.7 MB/s eta 0:00:00
```

```
Requirement already satisfied: numpy<2.0,>=1.17.3 in /usr/local/lib/python3.10/dist-pack
```

```
Requirement already satisfied: scipy>=1.5.0 in /usr/local/lib/python3.10/dist-packages (
```

```
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages
```

```
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-pa
```

```
Installing collected packages: scikit-learn
```

```
  Attempting uninstall: scikit-learn
```

```
    Found existing installation: scikit-learn 1.2.2
```

```
  Uninstalling scikit-learn-1.2.2:
```

```
    Successfully uninstalled scikit-learn-1.2.2
```

```
Successfully installed scikit-learn-1.3.2
```

```
WARNING: The following packages were previously imported in this runtime:
```

```
[sklearn]
```

```
You must restart the runtime in order to use newly installed versions.
```

RESTART RUNTIME

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
from sklearn.metrics import confusion_matrix
```

```
# Replace 'best_xgb' with your chosen model
```

```
y_pred = best_xgb.predict(X_valid)
```

```
cm = confusion_matrix(Y_valid, y_pred)
```

```
plt.figure(figsize=(8, 6))
```

```
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=["Predicted 0", "Predicted 1"], yticklabels=["Actual
```

```
plt.xlabel("Predicted")
```

```
plt.ylabel("Actual")
```

```
plt.title("Confusion Matrix")
```

```
plt.show()
```

